

Classification Level: Top Secret() Secret() Internal() Public(√)

RKLLM SDK User Guide

(Graphic Computing Platform Center)

Mark:	Version:	1.3.0
[] Changing	Author:	HPC
[√] Released	Completed Date:	2026-06-16
	Reviewer:	Vincent
	Reviewed Date:	2026-06-16

瑞芯微电子股份有限公司

Rockchip Electronics Co., Ltd.

(Copyright Reserved)

Revision History

Version	Modifier	Date	Modify description	Reviewer
V1.0.0	AI Group	2024-3-23	Initial version	Vincent
V1.0.1	AI Group	2024-5-8	1. Optimize memory usage. 2. Optimize inference time. 3. Optimize quantization accuracy. 4. Support Gemma, Phi-3 and other models. 5. Add Server call and interrupt interface.	Vincent
V1.0.1	AI Group	2024-10-10	1. Support group-wise quantization. 2. Support joint inference with lora model loading. 3. Support storage and preloading of prompt cache. 4. Optimize initialization, prefill, and decode time. 5. Support gguf model conversion. 6. Add gdq algorithm to improve 4-bit quantization accuracy. 7. Add mixed quantization algorithm, supporting a combination of grouped and non-grouped quantization based on specified ratios. 8. Add support for models such as Llama3, Gemma2, and Minicpm3.	Vincent
V1.1.4	AI Group	2025-2-7	1. Support for GPTQ model conversion 2. Added support for TeleChat2 and MiniCPM-S models 3. Added update_rkllm interface 4. Support for LLM model conversion in MiniCPMV and Qwen2VL 5. Fixed inference issues in MiniCPM3 6. Added DeepSeek_R1_Distill_Demo 7. Added Qwen2_VL_2B_Demo	Vincent
			1. Supports custom model conversion. 2. Supports chat_template configuration. 3. Enables multi-turn dialogue interactions.	

Version	Modifier	Date	Modify description	Reviewer
V1.2.0	AI Group	2025-4-8	4. Implements automatic prompt cache reuse for improved inference efficiency. 5. Expands maximum context length to 16K. 6. Supports embedding flash storage to reduce memory usage. 7. Introduces the GRQ Int4 quantization algorithm. 8. Supports GPTQ-Int8 model conversion. 9. Compatible with the RK3562 platform. 10. Added support for visual multimodal models such as InternVL2, Janus, and Qwen2.5-VL. 11. Supports CPU core configuration. 12. Added support for Gemma3 13. Added support for Python 3.9/3.11/3.12	Vincent
V1.2.1	AI Group	2025-6-25	1. Added support for RWKV7, Qwen3, and MiniCPM4 models 2. Added support for the RV1126B platform 3. Enabled function calling capability 4. Enabled cross-attention inference 5. Optimize the callback function to support pausing inference 6. Supported multi-batch inference 7. Optimized KV cache clearing interface 8. Improved chat template parsing with support for thinking mode selection 9. Server demo updated to support OpenAI-compatible format 10. Added return of model inference performance statistics 11. Supported mrope multimodal position encoding 12. A new quantization optimization algorithm has been added to improve quantization accuracy.	Vincent

Version	Modifier	Date	Modify description	Reviewer
V1.2.2	HPC	2025-9-29	1. Added support for Gemma3n and InternVL3 models 2. Supported for multi-instance inference 3. Supported for LongRoPE 4. Fixed issues with asynchronous inference interfaces 5. Fixed chat template parsing issues 6. Optimized inference performance 7. Optimized multimodal vision model demo	Vincent
V1.2.3	HPC	2025-11-22	1. Added support for InternVL3.5, DeepSeekOCR, and Qwen3-VL models 2. Added automatic cache reuse for embedding input 3. Added embedding input support for the Gemma3n model 4. Added support for loading chat template from an external file	Vincent
V1.3.0	HPC	2026-6-16	1. Added support for Qwen3.5, Gemma4, and SmoLM3 models. 2. Optimized the multimodal input interface and cache reuse strategy. 3. Added support for multiple EOS Token IDs and the ignore_eos_token parameter. 4. Optimized performance on 32-bit systems. 5. Added tokenizer callback and embedding callback interfaces. 6. Optimized long-context decoding performance for selected models on the RK3576 platform. 7. Improved the quantization method for embedding inputs. 8. Fixed incorrect memory usage statistics on the RV1126B platform. 9. Fixed numerical overflow issues during inference for selected models on the RK3588	Vincent

Version	Modifier	Date	Modify description	Reviewer
			platform. 10. Enhanced rkllm_server_demo compatibility with OpenAI API interfaces. 11. Added support for overriding max_new_tokens and sampling parameters in RKLLMInferParam.	

Rockchip

Table of contents

1 Introduction to RKLLM	8
1.1 Introduction to RKLLM Toolchain	8
1.1.1 Introduction to RKLLM-Toolkit	8
1.1.2 Introduction to RKLLM Runtime	8
1.2 Introduction to the RKLLM Development Process	9
1.3 Applicable Hardware Platforms	10
2 Development Environment Preparation	11
2.1 Installation of RKLLM-Toolkit	12
2.1.1 Installation via Pip	12
2.2 Introduction for RKLLM Runtime Usage	13
2.3 Compilation Requirements for RKLLM Runtime	13
2.4 Kernel Update	14
3 RKLLM Instruction Guide	16
3.1 Model Conversion	16
3.1.1 RKLLM Initialization	16
3.1.2 Loading Models	16
3.1.3 Construction for RKLLM Model	18
3.1.4 Exporting Model	21
3.1.5 GPTQ Model Conversion	22
3.1.6 Custom Model Conversion	23
3.1.7 Update For Old Version Model	26
3.1.8 Simulation Accuracy Evaluation	27
3.1.9 Simulated Model Inference	28
3.2 Inference Implementation in Board-side	29
3.2.1 Define Callback Function	30

3.2.2 Define RKLLMParam	32
3.2.3 Define Input Struct	35
3.2.4 Initialize Model	41
3.2.5 Inference Model	42
3.2.6 Interrupt Model Inference	42
3.2.7 Release Model	43
3.2.8 Load LoRA Model	43
3.2.9 Load Prompt Cache	44
3.2.10 KV Cache Management	46
3.2.11 Chat Template Settings	47
3.2.12 Function Calling Settings	48
3.2.13 Cross Attention Settings	52
3.2.14 Multi-batch Parallel Inference	53
3.2.15 Model Inference Pause	55
3.2.16 Model Performance Data Callback	55
3.2.17 tokenizer callback	56
3.2.18 embed callback	57
3.3 Board-side Inference Example	58
3.3.1 Complete Code of Example Project	58
3.3.2 Instructions for Example Project	59
3.3.3 Monitor inference performance and Log Viewing	59
3.4 Implementation of Board-side Server	60
3.4.1 Deployment Example of RKLLM-Server-Flask	61
3.4.2 Deployment Example of RKLLM-Server-Gradio	64
4 Reference	70

1 Introduction to RKLLM

1.1 Introduction to RKLLM Toolchain

1.1.1 Introduction to RKLLM-Toolkit

RKLLM-Toolkit is a development kit that provides users with the ability to quantify and convert large language models on the PC. Through the Python interface provided by the tool, the following functions can be conveniently accomplished:

1) Model Conversion: It supports the conversion of Large Language Model (LLM) in Hugging Face or GGUF format to RKLLM model, and the supported models include LLaMA, Qwen, Qwen2, Qwen3, Qwen2-VL, Qwen3-VL, Phi-2, Phi-3, ChatGLM3, Gemma, Gemma2, Gemma3, Gemma3n, Gemma4, InternLM2, TeleChat2, MiniCPM-S, MiniCPM, MiniCPM3, MiniCPM4, DeepSeekOCR, InternVL3, Qwen3.5, SmolLM3. The converted RKLLM models can be loaded and used on Rockchip NPU platform.

2) Quantization Function: Supports quantization of floating-point models into fixed-point models. Currently supported quantization types include:

- a. w4a16;
- b. w4a16 grouped quantization (supported group sizes: 32, 64, 128);
- c. w8a8;
- d. w8a8 grouped quantization (supported group sizes: 128, 256, 512);

1.1.2 Introduction to RKLLM Runtime

The RKLLM Runtime is primarily responsible for loading the RKLLM models obtained from the RKLLM-Toolkit conversion and executing inference via the Rockchip NPU. Throughout the RKLLM model inference process, users retain the autonomy to specify inference parameter configurations, define various text generation methodologies, and seamlessly retrieve inference results from the model via pre-defined callback functions.

1.2 Introduction to the RKLLM Development Process

The overall development process of RKLLM primarily comprises two phases: model conversion in PC and deployment and execution in board-side:

1) Model Conversion in PC:

During this phase, the provided Hugging Face formatted LLM is converted into RKLLM format for efficient inference on the Rockchip NPU platform. This step includes:

- a. Model Acquisition: Obtain the large language model from:
 - a) Open-source large language models in Hugging Face format.
 - b) Self-trained large language models, with the requirement that the saved model structure is consistent with models on the Hugging Face platform.
 - c) GGUF models, currently supporting only q4_0 and fp16 types.
- b. Model Loading: Load the model in huggingface format with the `rkllm.load_huggingface()` function and load the gguf format model with `rkllm.load_gguf()` function;
- c. Model Quantization Configuration: Utilize the `rkllm.build()` function to construct the RKLLM model, with the option to perform model quantization for enhanced hardware deployment performance. Additionally, choose different optimization levels and quantization types during the construction process.
- d. Model Export: Export the RKLLM model as a `.rkllm` format file using the `rkllm.export_rkllm()` function for subsequent deployment.

2) Deployment and Execution in Board-side:

This phase encompasses the actual deployment and execution of the model, typically involving the following steps:

- a. Model Initialization: Load the RKLLM model onto the Rockchip NPU platform, configure model parameters accordingly to define the desired text generation methods, and pre-define callback functions to receive real-time inference results, preparing for inference.
- b. Model Inference: Execute the inference operation by passing input data to the model and

running model inference. Users can continuously retrieve inference results through pre-defined callback functions.

c. Model Release: After completing the inference process, release the model resources to allow other tasks to utilize the computational resources of the NPU.

These two steps constitute the complete RKLLM development process, ensuring successful conversion, debugging, and ultimately efficient deployment of the LLMs on the Rockchip NPU.

1.3 Applicable Hardware Platforms

The hardware platforms to which this document applies mainly include: RK3576, RK3588, RK3562, RV1126B.

2 Development Environment Preparation

The RKLLM toolchain zip file contains the whl installer for the RKLLM-Toolkit, the associated files of the RKLLM Runtime library and reference sample code. The specific folder structure is shown below:

```
doc
├── Rockchip_RKLLM_SDK_CN.pdf      # RKLLM SDK Documentation (Chinese)
├── Rockchip_RKLLM_SDK_EN.pdf      # RKLLM SDK Documentation (English)

examples
├── multimodal_model_demo          # Multimodal Inference Example
├── rkllm_api_demo                 # Board-side Inference Example
└── rkllm_server_demo              # RKLLM-Server Deployment Example

rkllm-runtime
├── Android
│   ├── librkllm_api
│   │   ├── arm64-v8a
│   │   │   ├── librkllmrt.so      # 64 位 Runtime 库
│   │   │   └── armeabi-v7a
│   │   │       ├── librkllmrt.so  # 32 位 Runtime 库
│   │   │       └── include
│   │   └── rkllm.h
│   └── Linux
│       ├── librkllm_api
│       │   ├── aarch64
│       │   │   ├── librkllmrt.so  # 64 位 Runtime 库
│       │   │   └── armhf
│       │   │       ├── librkllmrt.so # 32 位 Runtime 库
│       │   │       └── include
│       └── rkllm.h

rkllm-toolkit
├── packages
│   └── rkllm_toolkit-x.x.x-cp3xx-cp3xx-linux_x86_64.whl
└── examples

rknpu-driver
└── rknpu_driver_x.x.x_XXXXXXX.tar.bz2

scripts
├── fix_freq_rk3576.sh             # Fixed-frequency Script for RK3576
├── fix_freq_rk3588.sh             # Fixed-frequency Script for RK3588
├── fix_freq_rk3562.sh             # Fixed-frequency Script for RK3562
└── fix_freq_rv1126b.sh           # Fixed-frequency Script for RV1126B
```

This chapter provides detailed instructions for installing the RKLLM-Toolkit and RKLLM Runtime. For specific usage instructions, please refer to the instructions in Chapter 3.

2.1 Installation of RKLLM-Toolkit

This section mainly describes how to install the RKLLM-Toolkit with pip install command. Users can refer to the following detailed instructions to complete the installation of the RKLLM-Toolkit toolchain.

2.1.1 Installation via Pip

2.1.1.1 Installation of Miniforge3

To avoid the need for multiple versions of Python environments, it is recommended to use miniforge3 to manage your Python environments.

Check if miniforge3 is installed and the version information of conda. If it is already installed, you can skip this section.

```
conda -V
# If "conda: command not found" is displayed, it indicates that conda
is not installed.
# For example, version information could be displayed as conda 23.9.0
```

Download the miniforge3 installation package.

```
wget -c https://mirrors.bfsu.edu.cn/github-release/conda-
forge/miniforge/LatestRelease/Miniforge3-Linux-x86_64.sh
```

Install miniforge3.

```
chmod 777 Miniforge3-Linux-x86_64.sh
bash Miniforge3-Linux-x86_64.sh
```

2.1.1.2 Create RKLLM-Toolkit Conda Environment

Activate the Conda base environment.

```
source ~/miniforge3/bin/activate # directory of miniforge3
# (base) xxx@xxx-pc:~$
```

Create a Conda environment named RKLLM-Toolkit with Python version 3.8 (recommended).

```
conda create -n RKLLM-Toolkit python=3.8
```

Activate the RKLLM-Toolkit Conda environment.

```
conda activate RKLLM-Toolkit
# (RKLLM-Toolkit) xxx@xxx-pc:~$
```

2.1.1.3 Installing RKLLM-Toolkit

In the RKLLM-Toolkit Conda environment, use the pip command to install the provided

toolchain .whl package directly. During the installation process, the installer will automatically download the necessary dependencies for the RKLLM-Toolkit tools.

```
pip3 install rkllm_toolkit-x.x.x-cp3xx-cp3xx-linux_x86_64.whl
```

If executing the following command does not result in any errors, the installation is successful.

```
python  
from rkllm.api import RKLLM
```

2.2 Introduction for RKLLM Runtime Usage

Within the RKLLM toolchain files provided, the following components are included for the RKLLM runtime:

- 1) librklmrt.so: RKLLM Runtime library suitable.
- 2) include/rkllm.h: Corresponding header file to librklmrt.so, containing descriptions of related structures and function definitions.

When constructing deployment inference code using the RKLLM toolchain, it is critical to ensure proper linkage to the above header file and function library to ensure compilation correctness. During the actual runtime of the code, it is equally important to ensure successful transfer of the above function library files to the board and complete library declaration through the following environment variable settings:

```
export LD_LIBRARY_PATH=/path/to/your/lib
```

2.3 Compilation Requirements for RKLLM Runtime

When utilizing RKLLM Runtime, it's essential to pay attention to the version of the gcc compilation tool. We recommend the cross-compilation tool [gcc-arm-10.2-2020.11-x86_64-aarch64-none-linux-gnu](#). Please note that cross-compilation tools often have backward compatibility but not upward compatibility, so refrain from using versions below 10.2.

If you choose to use the Android platform, and you need to compile Android executable files, we recommend using the Android NDK tool for cross-compilation. You can download it from the following link: [Android NDK Cross-Compilation Tool Download Link](#). We recommend using version r21e.

You can also refer to the specific compilation methods in the compilation scripts located in the

examples/rkllm_api_demo/deploy directory.

2.4 Kernel Update

Due to the requirement of a higher version of the NPU kernel for the provided RKLLM, users need to verify that the NPU kernel on the board is version v0.9.8 before using RKLLM Runtime for model inference. The specific query command is as follows:

```
# Execute the following command to query the NPU kernel version.
cat /sys/kernel/debug/rknpu/version

# Confirm that the command output is as follows:
# RKNPU driver: v0.9.8
```

If the queried NPU kernel version is lower than v0.9.8, please proceed to the official firmware address to download the latest firmware for updating.

For users using non-official firmware, it's necessary to update the kernel. The RKNPU driver package supports two main kernel versions: kernel-5.10 and kernel-6.1. For kernel-5.10, it is recommended to use a specific version at [GitHub-rockchip-linux/kernelatdevelop-5.10](#), and 5.10.209 is recommended. For kernel 6.1, it is recommended to use a specific version such as 6.1.84. Users can confirm the specific version number in the Makefile in the kernel's root directory. The specific steps to upgrade the kernel are as follows:

- 1) Download the [rknpu_driver_0.9.8_20241009.tar.bz2](#).
- 2) Unzip the compressed file and overwrite the RKNPU driver into the current kernel code directory(kernel/drivers/rknpu).
- 3) Recompile the kernel.
- 4) Flash the newly compiled kernel to the device.

If an error log, as shown in Figure 2-1, is encountered during the kernel compilation process:

```
drivers/rknpu/rknpu_gem.c: In function 'rknpu_gem_mmap_pages':
drivers/rknpu/rknpu_gem.c:891:2: error: implicit declaration of function 'vm_flags_set' [-Werror=implicit-function-declaration]
891 | vm_flags_set(vma, VM_MIXEDMAP);
    | ~~~~~
drivers/rknpu/rknpu_gem.c: In function 'rknpu_gem_mmap_buffer':
drivers/rknpu/rknpu_gem.c:988:2: error: implicit declaration of function 'vm_flags_clear' [-Werror=implicit-function-declaration]
988 | vm_flags_clear(vma, VM_PFNMAP);
    | ~~~~~
cc1: all warnings being treated as errors
make[2]: *** [scripts/Makefile.build:273: drivers/rknpu/rknpu_gem.o] Error 1
make[1]: *** [scripts/Makefile.build:516: drivers/rknpu] Error 2
make[1]: *** Waiting for unfinished jobs....
AR drivers/iio/built-in.a
make: *** [Makefile:1929: drivers] Error 2
MAKE KERNEL IMAGE FAILED.
```

Figure 2-1: Example of Kernel Compilation Error

Then need to modify the kernel header file kernel/include/linux/mm.h by adding the following code:

```
static inline void vm_flags_set(struct vm_area_struct *vma,
                               vm_flags_t flags){
    vma->vm_flags |= flags;
}

static inline void vm_flags_clear(struct vm_area_struct *vma,
                                  vm_flags_t flags){
    vma->vm_flags &= ~flags;
}
```

3 RKLLM Instruction Guide

3.1 Model Conversion

RKLLM-Toolkit provides model conversion and quantization functionalities. As one of the core features of RKLLM-Toolkit, it allows users to convert LLMs in Hugging Face format or GGUF format into RKLLM models, enabling the deployment and execution of RKLLM models on Rockchip NPU. This section will focus on the specific implementation of model conversion by RKLLM-Toolkit for LLMs, serving as a reference for users.

3.1.1 RKLLM Initialization

In this section, users need to initialize the RKLLM object, which is the first step in the entire workflow. In the example code, use the RKLLM() constructor function to initialize the RKLLM object.

3.1.2 Loading Models

After initializing RKLLM, users need to call the `rkllm.load_huggingface()` function to pass the specific path of the model. RKLLM-Toolkit will then successfully load the large language model in Hugging Face format or GGUF format based on the corresponding path, enabling subsequent conversion and quantization operations. The specific function definition is as follows:

Table 3-1 Interface Specification for the load_huggingface Function

Function	load_huggingface
Introduction	Used to load open-source LLMs in Hugging Face format.
Parameters	model: The path where the LLM files are stored, used for loading the model for subsequent conversion and quantization. model_lora: The file path of the LoRA weights. when converting, the model must point to the corresponding base model path. device: Specifies the device to be used for model conversion, supporting two options: cuda and cpu. dtype: The data type of the weights. Available options include float32, float16, and bfloat16. Using float16 or bfloat16 can reduce memory consumption but may affect quantization accuracy. custom_config: Custom configuration file for the model. For details, refer to the Custom Model Conversion section. load_weight: Whether to load the weights. If set to False, the actual weights will not be loaded.
Returns	0 indicates successful model loading; -1 indicates model loading failure.

The example code is as follows:

```
ret = rkllm.load_huggingface(  
    model = './huggingface_model_dir',  
    model_lora = './huggingface_lora_model_dir')  
if ret != 0:  
    print('Load model failed!')
```

Table 3-2 Interface Specification for the load_gguf Function

Function	load_gguf
Introduction	Used to load open-source large language models in GGUF format, supporting the numerical types q4_0 and fp16. GGUF-format LoRA models can also be loaded and converted to RKLLM models through this interface.
Parameters	model: The path where the LLM files are stored, used for loading the model for subsequent conversion and quantization.
Returns	0 indicates successful model loading; -1 indicates model loading failure.

The example code is as follows:

```
ret = rkllm.load_gguf(model = './model-Q4_0.gguf')
if ret != 0:
    print('Load model failed!')
```

3.1.3 Construction for RKLLM Model

After loading the original model through the rkllm.load_huggingface() function, the next step is to build the RKLLM model using the rkllm.build() function. During model conversation, users can choose whether to perform quantization, which helps reduce the model size and improve inference performance on Rockchip NPU. The specific definition of the rkllm.build() function is as follows:

Table 3-3 Interface Specification for the build Function

Function	build
Introduction	Used to construct the RKLLM model and define specific quantization operations during the conversion process.
Parameters	<i>do_quantization</i>: This parameter controls whether to perform quantization operations on the model, and it is recommended to set it to True. <i>optimization_level</i>: This parameter is used to set whether to perform quantization precision optimization. The available settings are {0, 1}, where 0 means no optimization and 1 means precision optimization. Precision optimization may cause a decrease in model inference performance. <i>quantized_dtype</i>: This parameter is used to set the specific type of quantization. Currently supported types include "w4a16", "w4a16_g32", "w4a16_g64", "w4a16_g128", "w8a8", "w8a8_g128", "w8a8_g256", "w8a8_g512". "w4a16" means 4-bit quantization for weights and no quantization for activations. "w4a16_g64" means 4-bit grouped quantization for weights (group size=64) and no quantization for activations. "w8a8" means 8-bit quantization for both weights and activations. "w8a8_g128" means 8-bit grouped quantization for both weights and activations (group size=128). The RK3576 and RV1126B platforms support five quantization types: "w4a16", "w4a16_g32", "w4a16_g64", "w4a16_g128", and "w8a8". The RK3588 platform supports four quantization types: "w8a8", "w8a8_g128", "w8a8_g256", and "w8a8_g512". The RK3562 platform supports those quantization types: "w8a8", "w4a16_g32", "w4a16_g64", "w4a16_g128", "w4a8_g32". For GGUF models, the quantization type corresponding to q4_0 is "w4a16_g32". Note: The group size must divide the output dimension of the linear layer; otherwise, quantization will fail! <i>quantized_algorithm</i>: Quantization accuracy optimization algorithm, with selectable

options being "normal" "grq". All quantization types can choose "normal", but the "grq" algorithm only support "w4a16" and grouped "w4a16" quantization, and it requires high computational power, with GPU acceleration being necessary. It supports improving quantization accuracy using LoRA modules. In the absence of any other imported LoRA models, you can use algorithms such as "normal_r[164]", or "grq_r[1~64]" — for example, "normal_r8", "grq_r8". In the end, both the quantized model and the LoRA model will be exported, and the runtime will load both models to perform inference.

num_npu_core: The number of NPU cores to be used for model inference. For "RK3576", the options are [1, 2], for "RK3562", the options are [1], for "RK3588", the options are [1, 2, 3] and for "RV1126B", the options are [1].

extra_qparams: Using the "grq" algorithm generate a *.qparams quantized weight cache file. This parameter can be set to the *.qparams path to rebegin the model export.

dataset: The dataset used for quantization calibration, formatted as JSON. An example of the content is as follows, where input is the question (with a prompt), and target is the answer. Multiple data entries are saved as dictionaries in a list;

hybrid_rate: The hybrid quantization rate ($\in [0,1)$). When the quantization type is "w4a16"/"w8a8", the model will mix "w4a16_g"/"w8a8_g" types at the specified rate to improve accuracy. When the quantization type is grouped "w4a16_g"/"w8a8_g", the model will mix "w4a16"/"w8a8" types at the specified rate to improve inference performance. When the hybrid_rate value is 0, no hybrid quantization will be performed.

target_platform: The hardware platform for running the model, with selectable options including "RK3576", "RK3588", "RK3562" and "RV1126B".

max_context: The maximum context length, supported up to 16,384 and must be

	aligned to 32.
Returns	0 indicates successful model conversion and quantization; -1 indicates model conversion failure.

The example code is as follows:

```
ret = rkllm.build(
    do_quantization=True,
    optimization_level=1,
    quantized_dtype='w8a8',
    quantized_algorithm="normal",
    num_npu_core=3,
    extra_qparams=None,
    dataset="quant_data.json",
    hybrid_rate=0,
    target_platform='rk3588')
```

3.1.4 Exporting Model

After constructing the RKLLM model using the `rkllm.build()` function, users can save the RKNN model as a `.rkllm` file for subsequent model deployment using the `rkllm.export_rkllm()` function. The specific parameter definition of the `rkllm.export_rkllm()` function is as follows:

Table 3-4 Interface Specification for the `export_rkllm` Function

Function	<code>export_rkllm</code>
Introduction	Used to save the converted and quantized RKLLM model for subsequent inference.
Parameters	export_path: The save path for exporting the RKLLM model file. LoRA models will automatically be saved with an <code>_lora</code> suffix in the RKLLM model filename. export_tokenizer: Whether to export the model tokenizer vocabulary. Default: True. export_embedding: Whether to export the model embedding layer. Default: True.
Returns	0 indicates successful export and saving of the model; -1 indicates model export failure.

The example code is as follows:

```
ret = rkllm.export_rkllm(export_path = './model.rkllm')
if ret != 0:
    print('Export model failed!')
```

3.1.5 GPTQ Model Conversion

The user can convert a floating-point model to a 4-bit/8-bit model using the [AutoGPTQ](#) open-source quantization tool before converting it to an RKLLM model, in addition to using the quantization algorithms provided in the tools above. When using AutoGPTQ to quantize the floating-point model, the following parameters must be configured:


```
bits=4
sym=true
group_size=32/64/128
desc_act=false
bits=8
sym=true
group_size=128/256/512
desc_act=false
```

Here's an example code for converting a GPTQ model in HuggingFace format to an RKLLM model:

```
modelpath = '/path/to/Model-Instruct-GPTQ-Int4'
llm = RKLLM()
ret = llm.load_huggingface(model=modelpath, model_lora = None,
device='cuda')
if ret != 0:
    print('Load model failed!')
    exit(ret)
# Build model
dataset = None
qparams = None
target_platform = "RK3576"
optimization_level = 1
quantized_dtype = "w4a16_g32" #w4a16_g64 or w4a16_g128
quantized_algorithm = "normal"
num_npu_core = 2
ret = llm.build(do_quantization=True,
optimization_level=optimization_level, quantized_dtype=quantized_dtype,
quantized_algorithm=quantized_algorithm,
target_platform=target_platform, num_npu_core=num_npu_core,
extra_qparams=qparams, dataset=dataset)
if ret != 0:
    print('Build model failed!')
    exit(ret)
# Export rkllm model
ret =
llm.export_rkllm(f"./{os.path.basename(modelpath)}_{quantized_dtype}_{
target_platform}.rkllm")
if ret != 0:
    print('Export model failed!')
    exit(ret)
```

3.1.6 Custom Model Conversion

If the user has modified the model structure or name, and the modified overall architecture is as follows, they can use the custom function to convert the model.



Taking the Qwen model as an example, the corresponding variable names in the modeling_qwen.py file should be filled into the custom configuration file as follows.

```
{
  "BLOCKNAME": "QWenBlock",
  "TOKEN_EMBD": "wte",
  "ATTN_NORM": "ln_1",
  "ATTN_Q_NORM": "",
  "ATTN_K_NORM": "",
  "CROSS_ATTN_NORM": "",
  "CROSS_ATTN_Q": "",
  "ATTN_Q": "",
  "ATTN_K": "",
  "ATTN_V": "",
  "ATTN_QKV": "attn.c_attn",
  "ATTN_KV": "",
  "KV_CONTINUOUS": "true",
  "ATTN_OUT": "attn.c_proj",
  "CROSS_ATTN_OUT": "",
  "ATTN_POST_NORM": "",
  "FFN_NORM": "ln_2",
  "FFN_UP": "mlp.w1",
  "FFN_GATE": "mlp.w2",
  "ACT_TYPE": "silu",
  "FFN_DOWN": "mlp.c_proj",
  "FFN_POST_NORM": "",
  "OUTPUT_NORM": "ln_f",
  "OUTPUT": "lm_head"
}
```

The available options for ACT_TYPE are ["silu", "gelu", "relu", "fatrelu", "squarerelu", "swiglu"].

ATTN_NORM and FFN_NORM only support RMSNorm.

If ATTN_QKV or ATTN_KV is used, it must be confirmed whether the weights can be split into contiguous K|V. If they are contiguous, set KV_CONTINUOUS to true. For example, the c_attn weights in the Qwen model are stored contiguously and can be sequentially split into Q, K, and V by size. In contrast, the wqkv weights in the InternLM2 model are stored non-contiguously, with Q, K, and V interleaved by head_dim.

The modeling_qwen.py file definition for the Qwen 1.8B model is as follows:


```

class QWenLMHeadModel(QWenPreTrainedModel):
    def __init__(self, config):
        super().__init__(config)
        self.transformer = QWenModel(config)
        self.lm_head = nn.Linear(config.hidden_size,
config.vocab_size, bias=False)
class QWenModel(QWenPreTrainedModel):
    _keys_to_ignore_on_load_missing = ["attn.masked_bias"]

    def __init__(self, config):
        super().__init__(config)
        self.wte = nn.Embedding(self.vocab_size, self.embed_dim)
        self.ln_f = RMSNorm(
            self.embed_dim,
            eps=config.layer_norm_epsilon,
        )
class QWenBlock(nn.Module):
    def __init__(self, config):
        super().__init__()
        hidden_size = config.hidden_size
        self.bf16 = config.bf16

        self.ln_1 = RMSNorm(
            hidden_size,
            eps=config.layer_norm_epsilon,
        )
        self.attn = QWenAttention(config)
        self.ln_2 = RMSNorm(
            hidden_size,
            eps=config.layer_norm_epsilon,
        )
        self.mlp = QWenMLP(config)
class QWenAttention(nn.Module):
    def __init__(self, config):
        super().__init__()

        self.c_attn = nn.Linear(config.hidden_size, 3 *
self.projection_size)

        self.c_proj = nn.Linear(
            config.hidden_size, self.projection_size, bias=not
config.no_bias
        )

class QWenMLP(nn.Module):
    def __init__(self, config):
        super().__init__()
        self.w1 = nn.Linear(
            config.hidden_size, config.intermediate_size // 2, bias=not
config.no_bias
        )
        self.w2 = nn.Linear(
            config.hidden_size, config.intermediate_size // 2, bias=not
config.no_bias
        )
        ff_dim_in = config.intermediate_size // 2
        self.c_proj = nn.Linear(ff_dim_in, config.hidden_size, bias=not
config.no_bias)

```

For custom model conversion, including models with cross-attention support, we provide a reference

Hugging Face model structure. The configuration file for customization is as follows:

```
{
  "BLOCKNAME": "CustomDecoderLayer",
  "TOKEN_EMBED": "embed_tokens",
  "ATTN_NORM": "input_layernorm",
  "ATTN_Q_NORM": "",
  "ATTN_K_NORM": "",
  "CROSS_ATTN_NORM": "cross_layernorm",
  "CROSS_ATTN_Q": "cross_attn.cross_q_proj",
  "ATTN_Q": "",
  "ATTN_K": "",
  "ATTN_V": "",
  "ATTN_QKV": "self_attn.qkv_proj",
  "ATTN_KV": "",
  "KV_CONTINUOUS": "true",
  "ATTN_OUT": "self_attn.o_proj",
  "CROSS_ATTN_OUT": "cross_attn.cross_o_proj",
  "ATTN_POST_NORM": "",
  "FFN_NORM": "post_attention_layernorm",
  "FFN_UP": "mlp.up_proj",
  "FFN_GATE": "mlp.gate_proj",
  "ACT_TYPE": "silu",
  "FFN_DOWN": "mlp.down_proj",
  "FFN_POST_NORM": "",
  "OUTPUT_NORM": "norm",
  "OUTPUT": "lm_head"
}
```

3.1.7 Update For Old Version Model

Since version 1.0.2 differs significantly from version 1.1 and later versions, the `rkllm.update_rkllm()` function is provided to update the 1.0.2 model to the latest version.. When updating the model, there is no need to perform the model loading and building steps mentioned earlier. Instead, the update can be done directly via this interface, and after the update, parameters such as model quantization type remain unchanged.

Here is a description of the function and its parameters:

Table 3-5 Interface Specification for the update_rkllm Function

Function	update_rkllm
Introduction	Update the model from version 1.0.2 to latest version
Parameters	model: The path to the RKLLM model of version 1.0.2.
Returns	0 indicates that the model update was successful; -1 indicates that the model update failed.

The example code is as follows:

```
ret = llm.update_rkllm(model = "./model_1.0.2version.rkllm")
if ret != 0:
    print('Load model failed!')
    exit(ret)
```

3.1.8 Simulation Accuracy Evaluation

After the user builds the RKLLM model through the rkllm.build() function, they can perform simulation accuracy evaluation on the PC using the rkllm.get_logits() function. The specific parameter definitions for the rkllm.get_logits() function are as follows:

Table 3-6 Interface Specification for the get_logits Function

Function	get_logits
Introduction	Used for simulation accuracy evaluation on the PC.
Parameters	inputs: The simulation input format is the same as Hugging Face model inference. An example is as follows: {"input_ids": "", "top_k": 1, ...}
Returns	The logits values inferred by the model.

Example code for using this function to perform PPL (perplexity) testing on the Wikitext dataset is as follows:

```
def eval_wikitext(llm):
    seqlen = 512
    tokenizer = AutoTokenizer.from_pretrained(
        modelpath,
        trust_remote_code=True
    )
    #Dataset download link:
    #https://huggingface.co/datasets/Salesforce/wikitext/tree/main/wikitext-2-raw-v1
    testenc = load_dataset("parquet", data_files='./wikitext/wikitext-2-raw-1/test-00000-of-00001.parquet', split='train')
    testenc = tokenizer(
        "\n\n".join(testenc['text']),
        return_tensors="pt").input_ids
    nsamples = testenc.numel() // seqlen
    nlls = []
    for i in tqdm(range(nsamples), desc="eval_wikitext: "):
        batch = testenc[:, (i * seqlen): ((i + 1) * seqlen)]
        inputs = {"input_ids": batch}
        lm_logits = llm.get_logits(inputs)
        if lm_logits is None:
            print("get logits failed!")
            return
        shift_logits = lm_logits[:, :-1, :]
        shift_labels = batch[:, 1:].to(lm_logits.device)
        loss_fct = nn.CrossEntropyLoss().to(lm_logits.device)
        loss = loss_fct(shift_logits.view(-1, shift_logits.size(-1)),
            shift_labels.view(-1))
        neg_log_likelihood = loss.float() * seqlen
        nlls.append(neg_log_likelihood)
    ppl = torch.exp(torch.stack(nlls).sum() / (nsamples * seqlen))
    print(f'wikitext-2-raw-1-test ppl: {round(ppl.item(), 2)}')
```

3.1.9 Simulated Model Inference

After the user builds the RKLLM model using the `rkllm.build()` function, they can perform simulated inference on the PC using the `rkllm.chat_model()` function. The specific parameter definitions for the `rkllm.chat_model()` function are as follows:

Table 3-7 Interface Specification for the `chat_model` Function

Function	<code>chat_model</code>
Introduction	Used for simulate model inference on the PC.
Parameters	messages: The text input, which needs to include the appropriate prompts. args: Inference configuration parameters, such as sampling parameters like <code>top_k</code> .
Returns	The logits values inferred by the model.

The example code is as follows:

```
args = {
    "max_length":128,
    "top_k":1,
    "temperature":0.8,
    "do_sample":True,
    "repetition_penalty":1.1
}

mesg = "Human: How's the weather today?\nAssistant:"
print(llm.chat_model(mesg, args))
```

The above operations cover all steps of model conversion and quantization in the RKLLM-Toolkit.

Depending on different requirements and application scenarios, users can choose different configuration options and quantization methods for customised settings, which facilitates subsequent deployment.

3.2 Inference Implementation in Board-side

This chapter introduces the usage of the general API interface functions. Users can refer to the content of this chapter to construct C++ code and implement inference of RKLLM models on the board to obtain inference results. The RKLLM board-side inference implementation is as follows:

- 1) Define the callback function RKLLMCallback.
- 2) Define the RKLLM model parameter structure RKLLMParam.
- 3) Initialize the RKLLM model with rkllm_init().
- 4) Perform model inference with rkllm_run().
- 5) Process the real-time inference results returned by the callback function result_callback.
- 6) Destroy the RKLLM model and release resources with rkllm_destroy().

In the subsequent parts of this chapter, the document will provide detailed explanations of each step in the process and provide detailed explanations of the functions involved.

3.2.1 Define Callback Function

The callback function is used to receive real-time output results from the RKLLM model. It is bound during the initialization of RKLLM and continuously outputs results to the callback function during the RKLLM model inference process, returning only one token each time.

Here is an example that callback function prints the output results in real-time to the terminal:

```
int callback(RKLLMResult* result, void* userdata, LLMCallState state)
{
    if(state == LLM_RUN_NORMAL){
        printf("%s", result->text);
    }
    if (state == LLM_RUN_FINISH) {
        printf("finish\n");
    } else if (state == LLM_RUN_ERROR){
        printf("\run error\n");
    }
    return 0;
}
```

1) LLMCallState is a status flag, and its specific definition is as follows:

Table 3-8 Explanation of LLMCallState Status Flags

Definition	LLMCallState
Introduction	Used to indicate the current running state of RKLLM.
Enumeration Values	LLM_RUN_NORMAL: indicates that the RKLLM model is currently inferencing; LLM_RUN_FINISH: indicates that the RKLLM model has completed inference; LLM_RUN_WAITING: indicates that the currently decoded character from RKLLM is not a complete UTF-8 encoding and needs to be concatenated with the next decoding result; LLM_RUN_ERROR: indicates that an error has occurred during inference;

During the design process of the callback function, users can set different post-processing behaviors based on the different states of LLMCallState.

2) RKLLMResult is a return value structure, and its specific definition is as follows:

Table 3-9 Explanation of RKLLMResult Structure

Definition	RKLLMResult
Introduction	Used to return the current inference-generated result.
Struct Fields	const char* text: indicates the text content generated by the current inference; int32_t token_id: indicates the token_id generated by the current inference; RKLLMResultLogits logits: Represents the logits information generated during the current inference, only returned when RKLLMInferMode is set to RKLLM_INFER_GET_LOGITS; RKLLMPerfStat perf: Indicates the performance statistics at the end of the current inference; data is only returned when in the RKLLM_RUN_FINISH state.

3) RKLLMResultLogits is a return value structure, and its specific definition is as follows:

Table 3-10 Explanation of RKLLMResultLogitsStructure

Definition	RKLLMResultLogits
Introduction	Used to return the current inference-generated logits.
Struct Fields	const float* logits: indicates the logits generated by the current inference; int vocab_size: indicates the token_id generated by the current inference; int num_tokens: Indicates the number of tokens returned; users can calculate the size of the logits based on vocab_size and num_tokens;

During the design process of the callback function, users can set different post-processing behaviors based on the values in RKLLMResult.

3.2.2 Define RKLLMParam

The RKLLMParam structure is used to describe and define the detailed information of RKLLM.

The specific definition is as follows:

Table 3-11 Explanation of RKLLMParam Structure

Definition	RKLLMParam
Introduction	Used to define the detailed parameters of the RKLLM model.
Struct Fields	<i>const char* model_path:</i> the path to the RKLLM model file; <i>int32_t max_context_len:</i> the maximum context length during inference; <i>int32_t max_new_tokens:</i> the maximum number of generated tokens in inferencing; <i>int32_t top_k:</i> top-k sampling is a text generation method that selects the next token only from the top-k tokens with the highest probabilities predicted by the model. This method helps reduce the risk of generating low-probability or meaningless tokens. A higher top-k value (e.g., 100) will consider more token choices, resulting in more diverse text generation, while a lower value (e.g., 10) will focus on the most probable tokens, generating more conservative text. The default value is 40; <i>float top_p:</i> top-p sampling, also known as nucleus sampling, is another text generation method that selects the next token from a group of tokens with cumulative probabilities of at least p. This method balances diversity and quality by considering the probabilities of tokens and the number of sampled tokens. A higher top-p value (e.g., 0.95) results in more diverse text generation, while a lower value (e.g., 0.5) generates more focused and conservative text. The default value is 0.9; <i>float temperature:</i> a hyperparameter that controls the randomness of generated text by adjusting the probability distribution of output tokens. A higher temperature (e.g., 1.5) makes the output more random and creative. When the temperature is high, the model considers more options with lower probabilities when selecting the next token, resulting in more diverse and unexpected outputs. A lower temperature (e.g., 0.5)

makes the output more focused and conservative. Lower temperatures mean that the model is more likely to choose high-probability tokens, resulting in more consistent and predictable outputs. In the extreme case of a temperature of 0, the model always chooses the most probable next token, resulting in identical outputs every time. To balance randomness and determinism and ensure that the output is neither overly uniform and predictable nor overly random and chaotic, the default value is 0.8;

float repeat_penalty: controls the occurrence of token sequence repetitions in the generated text, helping to prevent the model from generating repetitive or monotonous text. A higher value (e.g., 1.5) imposes a stronger penalty on repetitions, while a lower value (e.g., 0.9) is more lenient. The default value is 1.1;

float frequency_penalty: a factor for penalizing word/phrase repetition, reducing the probability of using words/phrases with higher frequencies overall and increasing the likelihood of using those with lower frequencies. This may lead to more diversified generated text, but could also result in text that is difficult to understand or not as expected. The range is [-2.0, 2.0], with a default value of 0;

float presence_penalty: reduces the likelihood of generating tokens that have already appeared in the generated text. The default value is 0;

int32_t mirostat: an algorithm actively maintaining the quality of generated text within the expected range during the text generation process. It aims to find a balance between coherence and diversity, avoiding low-quality output caused by excessive repetition (boredom trap) or incoherence (confusion trap). The values space is {0, 1, 2}, where 0 indicates not activating the algorithm, 1 indicates using the mirostat algorithm, and 2 indicates using the mirostat 2.0 algorithm;

float mirostat_tau: an option setting the target entropy for mirostat, representing the expected perplexity value of the generated text. Adjusting the target entropy allows to control the balance between coherence and diversity in the generated text. Lower

	values will result in more concentrated and coherent text, while higher values will lead to more diversified text, possibly with lower coherence. The default value is 5.0; <i>float mirostat_eta:</i> an option setting the learning rate for mirostat, which influences the algorithm's responsiveness to feedback on generated text. A lower learning rate will result in slower adjustment, while a higher learning rate will make the algorithm more sensitive. The default value is 0.1; <i>bool skip_special_token:</i> whether to skip special tokens and not output them, such as the end-of-sequence token <EOS>; <i>bool ignore_eos_token:</i> whether to ignore eos token; <i>bool is_async:</i> whether to use asynchronous mode; <i>n_keep:</i> The number of cache entries to retain at the beginning when clearing the KV cache. In multi-turn conversations, the n_keep value must be at least as long as the system_prompt length; <i>RKLLMExtendParam extend_param:</i> the special parameters for controlling inference.
--	---

Table 3-12 Explanation of RKLLMExtendParam Structure

Definition	RKLLMExtendParam
Introduction	The special parameters for controlling inference.
Struct Fields	int32_t base_domain_id: controls from which domain the RKLLM model starts initialization, default is 0. int8_t embed_flash: Controls whether to store the model's vocabulary in flash memory to save memory. Set to 0 to disable and 1 to enable. int8_t enabled_cpus_num: Sets the number of CPUs to use for inference. The range varies depending on the chip model. For RK3588/3576, the range is 1-8; for RK3562, the range is 1-4, with the default set to 4. uint32_t enabled_cpus_mask: Uses a binary mask to configure which specific CPU cores are used for inference. In rkllm.h, macros are predefined to represent CPU numbers, and you can configure them using the format CPU4 CPU5 CPU6 CPU7. uint8_t n_batch: Sets the level of parallel inference; the default value is 1. int8_t use_cross_attn: Sets whether to enable cross-attention.

In actual code construction, RKLLMParam needs to call the rkllm_createDefaultParam() function to initialize its definition, and set the corresponding model parameters according to requirements. Sample code is as follows:

```
RKLLMParam param = rkllm_createDefaultParam();
param.model_path = "model.rkllm";
param.top_k = 1;
param.max_new_tokens = 256;
param.max_context_len = 512;
```

3.2.3 Define Input Struct

To accommodate different input data types, the RKLLMInput input struct is defined, which currently accepts four types of input: text, image and text, token IDs, and encoded vectors. The specific definition is as follows:

Table 3-13 Explanation of RKLLMInput Structure

Definition	RKLLMInput
Introduction	Used to receive different forms of input data.
Struct Fields	<i>RKLLMInputType input_type</i>: Input mode; <i>const char* role</i>: Input type, optional values: ["user", "tool"]; <i>bool enable_thinking</i>: whether to enable the thinking mode of Qwen3; <i>union</i>: Used to store different input data types, specifically including the following forms: <i>const char prompt_input*</i>: Text prompt input, used to pass natural language text; <i>RKLLMEmbedInput embed_input</i>: Embedding vector input, representing processed feature vectors; <i>RKLLMTokenInput token_input</i>: Token input, used to pass the tokenized token sequence; <i>RKLLMMultiModalInput multimodal_input</i>: Multimodal input, which can pass multimodal data, such as combined input of images and text;

Table 3-14 Explanation of RKLLMInputType Structure

Definition	RKLLMInputType
Introduction	Used to represent the type of input data.
Enumeration Values	<i>RKLLM_INPUT_PROMPT</i>: Indicates that the input data is plain text. <i>RKLLM_INPUT_TOKEN</i>: Indicates that the input data is token IDs. <i>RKLLM_INPUT_EMBED</i>: Indicates that the input data is encoded vectors. <i>RKLLM_INPUT_MULTIMODAL</i>: Indicates that the input data consists of images and text.

When the input data is plain text, it can be directly input using `input_data`. When the input data consists of token IDs, encoded vectors, or images and text, the `RKLLMInput` must be used in conjunction with the `RKLLMTokenInput`, `RKLLMEmbedInput`, and `RKLLMMultiModalInput` structures. The

specific introduction is as follows:

1) RKLLMTokenInput is the input struct that receives token IDs. The specific definition is as follows:

Table 3-15 Explanation of RKLLMTokenInput Structure

Definition	RKLLMTokenInput
Introduction	Used to receive token_id data.
Struct Fields	<i>int32_t input_ids*</i> : Memory pointer for the input token IDs. <i>size_t n_tokens</i> : The number of tokens in the input data.

2) RKLLMEmbedInput is the input struct that receives encoded vectors. The specific definition is as follows:

Table 3-16 Explanation of RKLLMEmbedInput Structure

Definition	RKLLMEmbedInput
Introduction	Used to receive embedding data.
Struct Fields	<i>float embed*</i> : Memory pointer for the input token embeddings. <i>size_t n_tokens</i> : The number of tokens in the input data.

3) RKLLMMultiModalInput is the input struct that receives images and text. The specific definition is as follows:

Table 3-17 Explanation of RKLLMMultiModalInput Structure

Definition	RKLLMMultiModalInput
Introduction	Used to receive images and text data.
Struct Fields	<i>char prompt*</i> : Memory pointer for the input text. <i>float image_embed*</i> : Memory pointer for the input image embeddings. <i>size_t n_image_tokens</i> : The number of tokens for the input image embeddings. <i>size_t n_image</i> : number of input images; supports inputting multiple consecutive frames.

	<i>const char img_start*</i>: option to set the start marker for multimodal input image encoding, which needs to be configured in multimodal input mode. <i>const char img_end*</i>: option to set the end marker for multimodal input image encoding, which needs to be configured in multimodal input mode. <i>const char img_content*</i>: option to set the content marker for multimodal input image encoding, which needs to be configured in multimodal input mode. <i>size_t image_width</i>: width of the input image used when computing embeddings, used for mrope calculation parameters. <i>size_t image_height</i>: height of the input image used when computing embeddings, used for mrope calculation parameters.
	<i>float* video_embed</i>: memory pointer for the input video embeddings. <i>size_t n_frame_tokens</i>: number of tokens for each video frame. <i>size_t n_frame_per_video</i>: number of frames in each video. <i>size_t n_video</i>: number of videos, corresponding to the number of <video> tags in the prompt. <i>const char* video_start</i>: start tag of the video encoding. <i>const char* video_end</i>: end tag of the video encoding. <i>const char* video_content</i>: content tag of the video encoding. <i>size_t frame_width</i>: width of the video frame, used for MRoPE parameter calculation. <i>size_t frame_height</i>: height of the video frame, used for MRoPE parameter calculation.

RKLLM supports different inference modes and defines the RKLLMInferParam structure. It currently supports joint inference with preloaded LoRA models during the inference process, or saving a Prompt Cache for subsequent inference acceleration. The specific definition is as follows:

Table 3-18 Explanation of RKLLMInferParam Structure

Definition	RKLLMInferParam
Introduction	Used to define different inference modes.
Struct Fields	<i>RKLLMInferMode mode:</i> Inference mode, supporting RKLLM_INFER_GENERATE normal inference mode and RKLLM_INFER_GET_LOGITS additional logits retrieval inference mode. <i>RKLLMLoraParam* lora_params:</i> Parameter configuration for the LoRA used during inference, used to select which LoRA to infer when multiple LoRAs are loaded. Set to NULL if LoRA is not needed. <i>RKLLMPromptCacheParam* prompt_cache_params:</i> Parameter configuration for using the Prompt Cache during inference. Set to NULL if Prompt Cache generation is not needed. <i>RKLLMSamplingParam* sampling_params:</i> Pointer to sampling parameters used for model inference. If not specified, the sampling parameters configured during initialization will be used by default. <i>keep_history:</i> Indicates whether to retain the historical context during inference. Set to 1 for multi-turn conversations. <i>max_new_tokens:</i> Overrides the maximum number of tokens that can be generated in the current inference request.

Table 3-19 Explanation of RKLLMLoraParam Structure

Definition	RKLLMLoraParam
Introduction	Used to define the parameters for using LoRA during inference.
Struct Fields	<i>const char lora_adapter_name*:</i> The name of the LoRA used during inference.

Table 3-20 Explanation of RKLLMSamplingParam Structure

Definition	RKLLMSamplingParam
Introduction	Used to configure sampling parameters for model inference.
Introduction	int32_t top_k: Top-k sampling is a text generation method that selects the next token only from the top k tokens with the highest predicted probabilities. float top_p: Top-p sampling is a text generation method that selects the next token from the smallest set of tokens whose cumulative probability is at least p. float temperature: A hyperparameter that controls the randomness of text generation. Higher values increase diversity, while lower values make the output more deterministic. float repeat_penalty: Penalty factor applied to previously generated tokens to reduce repetition and help prevent repetitive or monotonous text generation. float frequency_penalty: Penalty factor based on token int32_t mirostat: Mirostat is an adaptive text generation algorithm that actively maintains the quality and perplexity of generated text within a desired range. float mirostat_tau: Target entropy parameter for Mirostat sampling. It represents the desired perplexity level of the generated text. float mirostat_eta: Learning rate parameter for Mirostat sampling. It controls how quickly the algorithm adapts to feedback from the generated text.

Table 3-21 Explanation of RKLLMPromptCacheParam Structure

Definition	RKLLMPromptCacheParam
Introduction	Used to define the parameters for using Prompt Cache during inference.
Struct Fields	int save_prompt_cache: Indicates whether to save the Prompt Cache during inference. 1 means it is required, and 0 means it is not. const char prompt_cache_path*: Path to save the Prompt Cache. If not set, it defaults to "./prompt_cache.bin".

Here is an example of using RKLLMPromptCacheParam for inference:

```
// Initialize and Set Prompt Cache Parameters(if needed)
RKLLMPromptCacheParam prompt_cache_params;
// Enable saving Prompt Cache
prompt_cache_params.save_prompt_cache = true;
// Specify cache file path
prompt_cache_params.prompt_cache_path = "./prompt_cache.bin";
rkllm_infer_params.prompt_cache_params = &prompt_cache_params;
```

3.2.4 Initialize Model

Before initializing the model, it is necessary to define the LLMHandle handle in advance. This handle is used for the initialization, inference, and resource release processes of the model. It's important to note that only by unifying the LLMHandle handle object across these three processes can the inference process of the model be completed correctly.

Prior to model inference, users need to complete the model initialization through the rkllm_init() function. The specific function definition is as follows:

Table 3-22 Interface Specification for the rkllm_init Function

Function	rkllm_init
Introduction	Used to initialize the specific parameters and inference settings for RKLLM model.
Parameters	LLMHandle* handle: register model to the corresponding handle for subsequent inference and release calls; RKLLMParam* param: the parameter structure defined for the model; RKLLMCallback*callback: pointer to the RKLLMCallback structure containing callback functions;
Returns	0 indicates that the initialization process is normal; -1 indicates initialization failure;

The example code is as follows:

```
LLMHandle llmHandle = nullptr;
RKLLMCallback rkllm_cb = {};
rkllm_cb.result_callback = callback;
rkllm_init(&llmHandle, &param, &rkllm_cb);
```

3.2.5 Inference Model

After completing the initialization process of the RKLLM model, users can perform model inference using the `rkllm_run()` function. Real-time inference results can be processed using the callback function predefined during initialization. The specific function definition of `rkllm_run()` is as follows:

Table 3-23 Interface Specification for the `rkllm_run` Function

Function	<code>rkllm_run</code>
Introduction	Used to performing result inference using the initialized RKLLM model.
Parameters	LLMHandle handle: the target handle registered during model initialization. RKLLMInput rkllm_input*: Input data for model inference. For details, see section 3.2.3 on input structure definition. RKLLMInferParam rkllm_infer_params*: Parameter passing during the model inference process. For details, see section 3.2.3 on input structure definition. void* userdata: user-defined function pointer. If set to NULL, <code>result_userdata</code> will be used.
Returns	0 indicates that the model inference runs normally; -1 indicates a failure in calling the model inference;

3.2.6 Interrupt Model Inference

During model inference, users can call the `rkllm_abort()` function to interrupt the inference process. The specific function definition is as follows:

Table 3-24 Interface Specification for the rkllm_abort Function

Function	rkllm_abort
Introduction	Used to interrupt the RKLLM model inference process.
Parameters	LLMHandle handle: the target handle registered during model initialization;
Returns	0 indicates successful interruption of the model; -1 indicates a failure to interrupt the model.

3.2.7 Release Model

After completing all model inference calls, users need to call the rkllm_destroy() function to destroy the RKLLM model and release the CPU, GPU, and NPU computing resources allocated, for use by other processes or models. The specific function definition is as follows:

Table 3-25 Interface Specification for the rkllm_destroy Function

Function	rkllm_destroy
Introduction	Used to destroy the RKLLM model and release all computing resources.
Parameters	LLMHandle handle: the target handle registered during model initialization;
Returns	0 indicates successful destruction and release of the RKLLM model; -1 indicates a failure in releasing the model.

The example code is as follows:

```
// llmHandle is the the target handle registered
rkllm_destroy(llmHandle);
```

3.2.8 Load LoRA Model

RKLLM supports running LoRA models simultaneously with the base model during inference. Before invoking the rkllm_run interface, you can load a LoRA model via the rkllm_load_lora interface. RKLLM allows loading multiple LoRA models; each call to rkllm_load_lora loads one LoRA model. The specific function definition is as follows:

Table 3-26 Interface Specification for the rkllm_load_lora Function

Function	rkllm_load_lora
Introduction	Used to load LoRA model for the base model.
Parameters	LLMHandle handle: The target handle registered during model initialization. See section 3.2.4 on initializing the model. RKLLMLoraAdapter lora_adapter*: Parameter for loading the LoRA model.
Returns	0 indicates the LoRA model was successfully loaded. -1 indicates the model loading failed.

Table 3-27 Explanation of RKLLMLoraAdapter Structure

Definition	RKLLMLoraAdapter
Introduction	Used to configure parameters when loading a LoRA model.
Struct Fields	const char* lora_adapter_path: The path to the LoRA model to be loaded. const char* lora_adapter_name: The name of the LoRA model to be loaded, defined by the user, used to select the specified LoRA during inference. float scale: The degree to which the LoRA model adjusts the base model parameters during inference.

Here is an example Code for Loading LoRA:

```
RKLLMLoraAdapter lora_adapter;
memset(&lora_adapter, 0, sizeof(RKLLMLoraAdapter));
lora_adapter.lora_adapter_path = "lora.rkllm";
lora_adapter.lora_adapter_name = "lora_name";
lora_adapter.scale = 1.0;
ret = rkllm_load_lora(llmHandle, &lora_adapter);
if (ret != 0) {
    printf("\nload lora failed\n");
}
```

3.2.9 Load Prompt Cache

During the model inference process, the Prefill stage typically consumes a significant amount of computational resources and time, especially when the Prompt is long. To accelerate this process,

RKLLM supports loading Prompt Cache from files. By reusing cached content, you can significantly reduce the time spent in the Prefill stage, thereby improving overall inference efficiency.

Before invoking the `rkllm_run` interface for inference, ensure that the `prompt_cache_params` parameters are correctly configured. This step allows the model to generate the corresponding Prompt Cache file after inference. When running inference for the first time, the system will automatically generate a Prompt Cache file. This file contains the intermediate results required for the Prefill stage, which can be reused in subsequent tasks. In subsequent inference tasks, you can load previously generated Prompt Cache files by calling the `rkllm_load_prompt_cache` interface.

The specific function definitions are as follows:

Table 3-28 Interface Specification for the `rkllm_load_prompt_cache` Function

Function	<code>rkllm_load_prompt_cache</code>
Introduction	Used to load Prompt Cache file.
Parameters	<i>LLMHandle handle:</i> The target handle registered during model initialization (refer to section 3.2.4 Initialization Model). <i>const char* prompt_cache_path:</i> The path to the Prompt Cache file to be loaded.
Returns	0 indicates the Prompt Cache file was successfully loaded. -1 indicates loading failed.

Table 3-29 Interface Specification for the `rkllm_release_prompt_cache` Function

Function	<code>rkllm_release_prompt_cache</code>
Introduction	Used to release the Prompt Cache.
Parameters	<i>LLMHandle handle:</i> The target handle registered during model initialization (refer to section 3.2.4 Initialization Model).
Returns	0 indicates that the Prompt Cache model was successfully released. -1 indicates that the model release failed.

Note:

RKLLM will detect the parts of the input that are the same as those in the prompt_cache from the beginning. If your input format is fixed as PROMPT_PREFIX + text + PROMPT_POSTFIX, you can generate the Prompt Cache for just the PROMPT_PREFIX part. After loading, you can reuse this part of the result in subsequent inferences.

RKLLM supports generating multiple Prompt Cache files. When different Prompt Cache files are needed, you can simply load the corresponding file. If you need to switch to another Prompt Cache file or no longer need the loaded Prompt Cache, please explicitly call the rkllm_release_prompt_cache interface to release it.

Here is an example Code for Loading Prompt Cache:

```
// Initialize and set the Prompt Cache parameters, then call the run
// interface to generate the Prompt Cache file.

RKLLMPromptCacheParam prompt_cache_params;
// Whether to save the prompt cache
prompt_cache_params.save_prompt_cache = true;
// If you need to save the prompt cache, specify the absolute path of
// the cache file.
prompt_cache_params.prompt_cache_path = "/data/prompt_cache.bin";
rkllm_infer_params.prompt_cache_params = &prompt_cache_params;
rkllm_infer_params.mode = RKLLM_INFER_GENERATE;
rkllm_input.input_type = RKLLM_INPUT_PROMPT;
rkllm_input.prompt_input = (char *)prompt.c_str();
rkllm_run(llmHandle, &rkllm_input, &rkllm_infer_params, NULL);
// Load the prompt cache file to reduce prefill time.
rkllm_load_prompt_cache(llmHandle, "./prompt_cache.bin");
if (ret != 0) {
    printf("\nload Prompt Cache failed\n");
}
rkllm_run(llmHandle, &rkllm_input, &rkllm_infer_params, NULL);
```

3.2.10 KV Cache Management

RKLLM supports manual clearing of the KV cache, which can be used for both single-turn and multi-turn dialogues. When invoking the cache clearing function, if keep_system_prompt is set to 1, the system prompt (if present) will be retained; otherwise, the entire cache will be cleared.

The function definition is as follows:

Table 3-30 Interface Specification for the rkllm_clear_kv_cache Function

Function	rkllm_clear_kv_cache
Introduction	Used to clear the KV cache.
Parameters	LLMHandle handle: The target handle registered during model initialization (refer to section 3.2.4 Initialization Model). keep_system_prompt: whether to retain the system prompt (1 to retain, 0 to clear). start_pos: The starting position ID of the KV cache to be deleted, inclusive of the start_pos position, If you don't need to use it, please manually set it to nullptr. end_pos: The ending position ID of the KV cache to be deleted, exclusive of the end_pos position, If you don't need to use it, please manually set it to nullptr. Note: deleting KV cache at specified positions is only supported when using the pause inference feature in single-turn mode. If a specific range [start_pos, end_pos) is provided, keep_system_prompt will be ignored, and the kv cache within this range will be cleared instead.
Returns	0: Indicates successful clearing of the KV cache. -1: Indicates failure to clear the KV cache.

3.2.11 Chat Template Settings

when users input text, RKLLM preprocesses the text. During preprocessing, it automatically parses and applies the prompt template based on the chat_template field in the HuggingFace model ' s tokenizer_config.json file. If customization is needed, you can reset it using the following function. Here, system_prompt acts as the system prompt to guide the model ' s behavior, prompt_prefix is the prefix added before the user input, and prompt_postfix is the suffix added after the user input. The specific function definition is as follows. After the user resets the template, the enable_thinking option becomes ineffective, and users need to configure it within their custom prompt.

Table 3-31 Interface Specification for the rkllm_set_chat_template Function

Function	rkllm_set_chat_template
Introduction	Used to set the prompt template.
Parameters	LLMHandle handle: The target handle registered during model initialization (refer to section 3.2.4 Initialization Model). system_prompt: The system prompt guiding the model behavior. prompt_prefix: The prefix before user input. prompt_postfix: The suffix after user input.
Returns	0: Indicates the success of the chat template. -1: Indicates failure to set the chat template.

3.2.12 Function Calling Settings

RKLLM supports structured interaction between the model and external systems through Function Calling, extending the capabilities of large language models and improving their performance in tasks such as knowledge augmentation and accurate data retrieval. when Function Calling mode is enabled, the application can pass function definitions to the model. Based on the user's query, the model determines whether a function call is needed and outputs a call request in the specified format. The application then calls the corresponding function based on the model's intent and returns the result. The model continues the conversation based on the returned result. the function used to configure Function Calling in RKLLM is defined as follows:

Table 3-32 Interface Specification for the rkllm_set_chat_template Function

Function	rkllm_set_function_tools
Introduction	Function Calling configuration, including system prompt, tool function definitions, and tool response identifier.
Parameters	LLMHandle handle: The target handle registered during model initialization (refer to section 3.2.4 Initialization Model). system_prompt: The system prompt guiding the model behavior. tools: json-formatted function definition string that describes the available functions, including their names, purposes, and parameter formats. tool_response_str: Identifier tag for tool function call results, used to distinguish them from regular conversation content.
Returns	0: indicates successful configuration; -1: indicates configuration failure.

Here is the example code, first, configure the tools:

```
std::string system_prompt = "You are Qwen, created by Alibaba Cloud.  
You are a helpful assistant.\n\nCurrent Date: 2024-09-30";  
std::string tools = R"(  
    {  
        "type": "function",  
        "function": {  
            "name": "get_current_temperature",  
            "description": "Get current temperature at a location.",  
            "parameters": {  
                "type": "object",  
                "properties": {  
                    "location": {  
                        "type": "string",  
                        "description": "The location to get the  
temperature for, in the format \"City, State, Country\"."  
                    },  
                    "unit": {  
                        "type": "string",  
                        "enum": ["celsius", "fahrenheit"],  
                        "description": "The unit to return the temperature  
in. Defaults to \"celsius\"."  
                    }  
                },  
                "required": ["location"]  
            }  
        }  
    }  
)
```

```

    },
    {
        "type": "function",
        "function": {
            "name": "get_temperature_date",
            "description": "Get temperature at a location and date.",
            "parameters": {
                "type": "object",
                "properties": {
                    "location": {
                        "type": "string",
                        "description": "The location to get the
temperature for, in the format \"City, State, Country\"."
                    },
                    "date": {
                        "type": "string",
                        "description": "The date to get the temperature
for, in the format \"Year-Month-Day\"."
                    },
                    "unit": {
                        "type": "string",
                        "enum": ["celsius", "fahrenheit"],
                        "description": "The unit to return the temperature
in. Defaults to \"celsius\"."
                    }
                },
                "required": ["location", "date"]
            }
        }
    }
])";

rkllm_set_function_tools(llmHandle, system_prompt.c_str(),
tools.c_str(), "tool_response");

```

Next, during the inference process, integrate with `rkllm_run` to implement the function calling chain:

```
// User's question
RKLLMInferParam rkllm_infer_params;
memset(&rkllm_infer_params, 0, sizeof(RKLLMInferParam));
rkllm_infer_params.mode = RKLLM_INFER_GENERATE;
rkllm_infer_params.keep_history = 0;

RKLLMInput rkllm_input;
rkllm_input.input_type = RKLLM_INPUT_PROMPT;
rkllm_input.enable_thinking = false;

rkllm_input.role = "user";
rkllm_input.prompt_input = "What's the temperature in San Francisco
now? How about tomorrow?";

rkllm_run(llmHandle, &rkllm_input, &rkllm_infer_params, NULL);

// the first call to rkllm_run will return the name of the tool
function that needs to be called;

<tool_call>
{"name": "get_current_temperature", "arguments": {"location": "San
Francisco"}}
</tool_call>
<tool_call>
{"name": "get_temperature_date", "arguments": {"location": "San
Francisco", "date": "2024-10-01"}}
</tool_call>

// return the tool call result to llm, and the role must be set to
"tool"

rkllm_input.role = "tool";
rkllm_input.prompt_input = R"([
    {
        "temperature": 26.1,
        "location": "San Francisco",
        "unit": "celsius"
    },
    {
        "temperature": 25.9,
        "location": "San Francisco",
        "date": "2024-09-30",
        "unit": "celsius"
    }
])";

rkllm_run(llmHandle, &rkllm_input, &rkllm_infer_params, NULL);

// final result
"The current temperature in San Francisco is 26.1 °C. Tomorrow, the
temperature is expected to be 25.9 °C."
```

3.2.13 Cross Attention Settings

RKLLM supports cross-attention inference. Users can input the K/V cache, mask, and positional information generated by the encoder into the decoder for cross-attention computation through the following function. Cross-attention is only supported for custom model inference. The model conversion method is detailed in Section 3.1.6.

Table 3-33 Interface Specification for the rkllm_set_cross_attn_params Function

Function	rkllm_set_cross_attn_params
Introduction	Used to configure cross-attention parameters.
Parameters	LLMHandle handle: The target handle registered during model initialization (refer to section 3.2.4 Initialization Model). RKLLMCrossAttnParam* cross_attn_params: Cross-attention parameters.
Returns	0: indicates successful configuration; -1: indicates configuration failure.

Table 3-34 Explanation of RKLLMCrossAttnParam Structure

Function	RKLLMCrossAttnParam
Introduction	Cross-attention parameters
Struct Fields	-float* encoder_k_cache: pointer to the encoder output key cache -float* encoder_v_cache: pointer to the encoder output value cache -float* encoder_mask: attention mask for the encoder -int32_t* encoder_pos: positional information of the input tokens to the encoder -int* num_tokens: number of input tokens to the encoder

3.2.14 Multi-batch Parallel Inference

RKLLM supports simultaneous inference of multiple batches (it is recommended that the number of batches does not exceed 8). Below is an example code for inference using two batches:

- When initializing the model, the `param.extend_param.n_batch` parameter needs to be set to 2.
- When using multiple batches for inference, both `RKLLMInput` and `RKLLMResult` in the callback are arrays of size `n_batch`.
- The inference results for all batches are returned synchronously in the callback. When the token ID of a particular batch is negative, it indicates that the inference for that batch is complete. Inference will only stop once all batches have finished.
- When handling the returned text in the callback, it is necessary to first check whether the returned text is a null pointer before performing assignment operations.

```
#include <string.h>
#include <unistd.h>
#include <string>
#include "rkllm.h"
#include <fstream>
#include <iostream>
#include <csignal>
#include <vector>

using namespace std;
LLMHandle llmHandle = nullptr;

std::string output_texts[10];

int callback(RKLLMResult *result, void *userdata, LLMCallState state)
{
    if (state == RKLLM_RUN_FINISH)
    {
        printf("\nrkllm run finish\n");
    } else if (state == RKLLM_RUN_ERROR) {
        printf("\nrkllm run error\n");
    } else if (state == RKLLM_RUN_NORMAL) {
        RKLLMResult batch1 = result[0];
        RKLLMResult batch2 = result[1];
        if (batch1.text) {
            output_texts[0] += batch1.text;
            printf("batch 0 %s\n", output_texts[0].c_str());
        }
        if (batch2.text) {
            output_texts[1] += batch2.text;
            printf("batch 1 %s\n", output_texts[1].c_str());
        }
    }
    return 0;
}

int main(int argc, char **argv)
{
    if (argc < 4) {
        std::cerr << "Usage: " << argv[0] << " model_path
max_new_tokens max_context_len\n";
        return 1;
    }

    RKLLMParam param = rkllm_createDefaultParam();
    param.model_path = argv[1];
    param.top_k = 1;
    param.top_p = 0.95;
    param.temperature = 0.8;
    param.repeat_penalty = 1.1;
    param.frequency_penalty = 0.0;
    param.presence_penalty = 0.0;
    param.max_new_tokens = std::atoi(argv[2]);
    param.max_context_len = std::atoi(argv[3]);
    param.skip_special_token = true;
    param.extend_param.base_domain_id = 0;
    param.extend_param.embed_flash = 1;
    param.extend_param.n_batch = 2;
```

```

RKLLMCallback rkllm_cb = {};
rkllm_cb.result_callback = callback;
int ret = rkllm_init(&llmHandle, &param, &rkllm_cb);
if (ret == 0){
    printf("rkllm init success\n");
} else {
    printf("rkllm init failed\n");
    exit_handler(-1);
}
RKLLMInput rkllm_input[2];
memset(&rkllm_input, 0, sizeof(RKLLMInput)*2);
RKLLMInferParam rkllm_infer_params;
memset(&rkllm_infer_params, 0, sizeof(RKLLMInferParam));
output_texts[0].clear();
output_texts[1].clear();

rkllm_infer_params.mode = RKLLM_INFER_GENERATE;
rkllm_infer_params.keep_history = 0;

rkllm_input[0].input_type = RKLLM_INPUT_PROMPT;
rkllm_input[0].role = "user";
rkllm_input[0].enable_thinking = false;
rkllm_input[0].prompt_input = "上联：江边惯看千帆过";

rkllm_input[1].input_type = RKLLM_INPUT_PROMPT;
rkllm_input[1].role = "user";
rkllm_input[1].enable_thinking = false;
rkllm_input[1].prompt_input = "以咏梅为题目，帮我写一首古诗，要求包含梅花、白雪等元素。";

printf("robot: ");
rkllm_run(llmHandle, &rkllm_input[0], &rkllm_infer_params, NULL);

rkllm_destroy(llmHandle);
return 0;
}

```

3.2.15 Model Inference Pause

RKLLM supports pausing inference in single-turn mode by returning 1 in the result callback. When inference is paused, the KV cache will not be cleared, allowing the user to modify the input and then call `rkllm_run` to resume inference.

3.2.16 Model Performance Data Callback

RKLLM supports returning performance data for a single inference at the end of the inference process. This includes the total time spent on prefill and decoding inference, the number of tokens, and memory usage. The structure of the returned data is defined as follows:

Table 3-35 Explanation of RKLLMPerfStat Structure

Definition	RKLLMPerfStat
Introduction	return model inference performance data
Struct Fields	- <i>float prefill_time_ms</i>: total time spent in the prefill stage, in milliseconds - <i>int prefill_tokens</i>: total number of tokens in the prefill stage - <i>float generate_time_ms</i>: total time spent in the generate stage, in milliseconds - <i>int generate_tokens</i>: total number of tokens in the generate stage - <i>float memory_usage_mb</i>: Memory usage during inference process, measured by VmHWM, in megabytes

3.2.17 tokenizer callback

Used to encode input text into a sequence of tokens, typically implemented by a tokenizer. The user needs to implement this callback function to convert a string (such as a prompt) into the required token ID format. Only effective when the model does not have a built-in tokenizer.

Table 3-36 Explanation of LLMTokenizerCallback Structure

Definition	LLMTokenizerCallback
Introduction	encode the specified text into a token sequence
Struct Fields	- <i>userdata</i>: user-defined parameter, typically a pointer to a tokenizer instance. - <i>text</i>: the input text to be encoded - <i>text_len</i>: length of text - <i>tokens</i>: array used to store the output token IDs - <i>n_tokens_max</i>: maximum number of tokens that the tokens array can hold

An example implementation is provided below:


```
int tokenizer_callback(void* userdata, const char* text, int32_t
text_len, int32_t* tokens, int32_t n_tokens_max)
{
    Tokenizer* tokenizer = (Tokenizer*)userdata;
    int n_tokens = tokenizer->Tokenize(text, text_len, tokens,
n_tokens_max);
    if (n_tokens <= 0) {
        printf("tokenizer failed for %s\n", text);
        return n_tokens;
    }
    return n_tokens;
}
RKLLMCallback rkllm_cb = {};
rkllm_cb.tokenizer_callback = tokenizer_callback;
rkllm_cb.tokenizer_userdata = tokenizer;
```

3.2.18 embed callback

Used to retrieve the embedding vectors corresponding to a sequence of tokens. In most cases, the user needs to implement this callback to look up the embedding for each input token and fill the embedding buffer..Only effective when the model does not have a built-in embedding layer.

Table 3-37 Explanation of LLMGetEmbedCallback Structure

Definition	LLMGetEmbedCallback
Introduction	Retrieves embedding data for the specified token sequence
Struct Fields	- userdata: user-defined parameter, typically a pointer to an embedding_info structure containing embedding-related information - tokens: array of input token IDs - num_tokens: number of tokens in the tokens array - embed: embedding buffer where the embedding vectors of all tokens should be stored sequentially - len: total size of the embedding buffer in bytes (should be equal to num_tokens * embedding_dim * sizeof(float32))

An example implementation is provided below:

```
struct embedding_info
{
    int      fd;
    float*   embedding_data;
    int      embedding_dim;
    int      vocab_size;
};

int embed_callback(void* userdata, int32_t* tokens, uint64_t
num_tokens, void* embed, uint64_t len)
{
    struct embedding_info* embed_info = (struct
embedding_info*)userdata;

    if (len != num_tokens * embed_info->embedding_dim * sizeof(float))
    {
        printf("num_tokens: %lu, embedding_dim: %d, expected embed
buffer size: %lu, got: %lu\n",
                num_tokens, embed_info->embedding_dim,
                num_tokens * embed_info->embedding_dim * sizeof(float),
len);
        return -1;
    }
    for (int n = 0; n < num_tokens; n++) {
        memcpy((unsigned char*)embed + n * embed_info->embedding_dim *
sizeof(float),
                embed_info->embedding_data + tokens[n] * embed_info-
>embedding_dim, embed_info->embedding_dim * sizeof(float));
    }

    return 0;
}

RKLLMCallback rkllm_cb = {};
rkllm_cb.embed_callback = embed_callback;
rkllm_cb.embed_userdata = &embed_info;
```

3.3 Board-side Inference Example

The directory, `examples/rkllm_api_demo`, includes the C++ project for the inference on the board-side, and which is synchronously updated with this documentation. Furthermore, compilation scripts are provided for users to facilitate the compilation of the project and the completion of the board-side inference of the RKLLM model.

3.3.1 Complete Code of Example Project

The complete C++ code example for inference calls is located in the `example/rkllm_api_demo/deploy/src` directory of the toolchain. The `llm_demo.cpp` file serves as an example of large language model inference. These examples include the full process, including model initialization, inference, handling outputs with callback functions, and releasing model resources. Users

can refer to the relevant code to implement custom functionality.

3.3.2 Instructions for Example Project

Under the directory of `examples/rkllm_api_demo/deploy`, not only does it contain sample code for invoking the RKLLM model inference, but also includes compilation scripts `build-android.sh` and `build-linux.sh`. This section will provide a brief explanation of how to use the sample code, taking compiling executable files for Linux systems as an example using the `build-linux.sh` script:

Firstly, users need to prepare cross-compilation tools on their own, noting that the recommended version of the compilation tools in section 2.3 is 10.2 or above. Subsequently, before the compilation process, users should replace the path to the cross-compilation tools in `build-linux.sh` themselves:

```
# Set the path of the cross-compiler GCC_COMPILER_PATH=~/.gcc-arm-10.2-2020.11-x86_64-aarch64-none-linux-gnu/bin/aarch64-none-linux-gnu
```

Subsequently, users can use `build-linux.sh` to initiate the compilation process. Upon completion of the compilation, users will obtain the corresponding `llm_demo` program, which will be installed in the `install/demo_linux_aarch64/llm_demo` directory. Following this, the executable file, library folder, and the RKLLM model (previously converted and quantized using the RKLLM-Toolkit tool) should be pushed to the board-side.

```
adb push install/demo_linux_aarch64 /data
adb push /PC/path/to/your/rkllm/model /data/demo_linux_aarch64
```

After completing the above steps, users can enter the terminal interface of the board using `adb`, and navigate to the corresponding `/data/demo_linux_aarch64` directory. Then, the inference of RKLLM model on the board can be invoked using the following command:

```
adb shell
cd /data/demo_linux_aarch64
export LD_LIBRARY_PATH=./lib
./llm_demo /path/to/your/rkllm/model 1024 2048
```

With the above operations, users can enter the example inference interface, interact with the board-side model for inference, and obtain real-time inference results from the RKLLM model.

3.3.3 Monitor inference performance and Log Viewing

If you need to view RKLLM's performance or logs during inference on the board, you can use the following commands:

```
# View only TTFT, TPS, and memory usage data
export RKLLM_LOG_LEVEL=1
# In addition to performance data, also prints cache length and more
details
export RKLLM_LOG_LEVEL=2
I rkllm: -----
I rkllm: Model init time (ms) 834.82
I rkllm: -----
I rkllm: Stage          Total Time (ms) Tokens    Time per Token (ms)    Tokens per Second
I rkllm: -----
I rkllm: Prefill         306.31      13      23.56      42.44
I rkllm: Generate        1918.31     9      213.15     4.69
I rkllm: -----
I rkllm: Peak Memory Usage (GB)
I rkllm: 0.52
I rkllm: -----
```

Figure 3-1 RKLLM inference performance logs on the hardware platform

3.4 Implementation of Board-side Server

After using RKLLM-Toolkit to convert the model and obtain the RKLLM model, users can deploy server-side services on Linux development boards. This involves setting up a server on a Linux device and exposing network interfaces to everyone in the local area network. Subsequently, the RKLLM model can be accessed by other users via the specified address, thus facilitating efficient and concise interaction. This section will introduce two different server deployment implementations.

1) RLM-Server-Flask based on Flask: Users can achieve API access between the client and server using request requests. The provided RKLLM-Server-Flask example includes a common conversation example and a single-round conversation example that supports the Function Calling function.

2) RKLLM-Server-Gradio based on Gradio: By reference to the provided example, users can rapidly construct a web server for visual interaction. Furthermore, the example illustrates the utilisation of the Gradio API interface, which enables users to undertake secondary development.

Both examples of server implementations mentioned above are located in the examples/rkllm_server_demo directory. This directory contains specific code for both implementations, one-click deployment scripts, and API interface calling examples. Users can choose different examples for reference and further development. The directory structure is as follows:

```
examples/rkllm_server_demo
├──rkllm_server          # Board-side Deployment Required Files
│   ├──lib              # RKLLM Runtime
│   ├──flask_server.py   # RKLLM-Server-Flask Example
│   └──gradio_server.py  # RKLLM-Server-Gradio Example
├──build_rkllm_server_flask.sh # One-click Deployment Script -Flask
├──build_rkllm_server_gradio.sh # One-click Deployment Script -Gradio
├──chat_api_flask.py      # API Interface Example -Flask
├──chat_api_gradio.py     # API Interface Example -Gradio
└──Readme.md
```

3.4.1 Deployment Example of RKLLM-Server-Flask

In the deployment example of RKLLM-Server-Flask, the main focus is on using the Flask framework to set up the server-side. On the client-side, data is transmitted using the request-response structure to achieve API access.

3.4.1.1 Server-side: One-click Deployment Script for RKLLM-Server-Flask

The one-click deployment script for RKLLM-Server-Flask is named `build_rkllm_server_flask.sh` and is located in the `rkllm_server_demo` directory. This script helps users quickly set up the RKLLM-Server-Flask server on a Linux development board. Before using this script, users should note the following:

- 1) Ensure that the development board is connected to the network via an Ethernet cable. Use the "ifconfig" command in the adb shell to determine the specific IP address of the development board. The RKLLM-Server-Flask will then set up the server with this IP address within the local area network and accept client access.

- 2) Users should have successfully converted the RKLLM model beforehand. Before executing the one-click deployment script, ensure that the RKLLM model has been pushed to the Linux board.

Users can directly invoke the `build_rkllm_server_flask.sh` script from their PC (not on a development board) to quickly deploy the RKLLM-Server-Flask server on a Linux development board.

The specific usage of the one-click deployment script `build_rkllm_server_flask.sh` is as follows:

```
./build_rkllm_server_flask.sh
  --workshop [RKLLM-Server Working Path]
  --model_path [Absolute Path of Converted RKLLM Model on Board] --
platform [Target Platform: RK3588/RK3576/RK3562/RV1126B]
  [--lora_model_path [LoRA Model Path]]
  [--prompt_cache_path [Prompt Cache File Path]] [--adb_device [ADB
Device Serial]]
```

The `workshop` parameter specifies the subsequent working directory of RKLLM-Server-Flask on the board. The `model\_path` parameter indicates the absolute path of the RKLLM model on the board, which was converted using RKLLM-Toolkit, and RKLLM-Server-Flask will read the model from this path during operation. The `platform` parameter specifies the current platform type, either RK3588, RK3576, RK3562 or RV1126B. The `lora\_model\_path` and `prompt\_cache\_path` parameters are optional and can be used to specify file paths when the user needs to load a LoRA model or use the prompt feature.

The following is a simple example of how to use the one-click deployment script `build\_rkllm\_server\_flask.sh`:

```
./build_rkllm_server_flask.sh
  --workshop /path/to/workshop
  --model_path /path/to/model.rkllm
  --platform rk3588 --adb_device [ADB Device Serial]
```

After executing above command, the one-click deployment script will perform the following steps:

- 1) Check the Linux environment on the board.
- 2) Automatically install the Flask library if not already installed.
- 3) Push the necessary files under rkllm_server_demo/rkllm_server to the board.
- 4) Index the RKLLM model in the preset working directory for RKLLM-Server-Flask.

Once you see the message "RKLLM Model has been initialized successfully!" in the terminal, it indicates that the RKLLM-Server-Flask example has been successfully launched.

```
=====init....=====
rkllm-runtime version: 1.0.2b9, rknpu driver version: 0.9.7, platform: RK3588
load prompt cache from '/data/cw/prompt_cache.bin'
loaded a prompt cache with prompt size of 27 tokens
RKLLM Model has been initialized successfully!
=====
* Serving Flask app 'flask_server'
* Debug mode: off
WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
* Running on all addresses (0.0.0.0)
* Running on http://127.0.0.1:8080
* Running on http://172.16.10.79:8080
Press CTRL+C to quit
```

Figure 3-2 Successful deployment of RKLLM-Server-Flask in terminal

By referring to the specific code logic in `build_rkllm_server_flask.sh`, users can understand the detailed deployment process of the RKLLM-Server-Flask example. This enables users to customize the deployment implementation of their server more flexibly. It is important to emphasize that in step 3 of the one-click deployment script, the script automatically synchronizes the current version of RKLLM Runtime to `rkllm_server/lib/librkllmrt.so`. This ensures that `flask_server.py` calls the current version of `librkllmrt.so` during runtime.

3.4.1.2 Server-side: Introductions for RKLLM-Server-Flask Example

In this section, we will outline and introduce the implementation approach of the deployment example for RKLLM-Server-Flask, helping users understand the construction logic of the example code for potential secondary development.

The deployment example of RKLLM-Server-Flask primarily relies on the Flask library to achieve the basic implementation of the server. Additionally, for RKLLM model inference, the `ctypes` library in Python is chosen to directly call the RKLLM Runtime library.

In the overall implementation of `rkllm_server/flask_server.py`, in order to call `librkllmrt.so` via `ctypes`, it's necessary to define relevant structures in Python based on the header file `rkllm.h` corresponding to `librkllmrt.so` beforehand. After the Flask server receives the struct data sent by users, it calls relevant functions to perform inference with the RKLLM model. The specific code implementation of `flask_server.py` mainly consists of the following steps:




Figure 3-3 Overview of the RKLLM-Server-Flask Deployment Implementation Process

- 1) Load RKLLM Runtime: Set the path for the dynamic library and load the dynamic library

librkllmrt.so using ctypes to achieve the analysis of the RKLLM Runtime.

2) Completing the external implementation of the Runtime: In the Python code, use ctypes to complete the external definitions of relevant implementations from the RKLLM Runtime header files, including definitions of data structures, callback functions, initialization functions, inference functions, destroy functions, and loading functions for LoRA/Prompt Cache.

3) Encapsulate a complete class function: Based on the various Runtime data types and function interfaces implemented in step 2), encapsulate complete class functions that integrate RKLLM's initialization, inference, destroy, and other operations for easier subsequent calls.

4) Set up RKLLM-Server-Flask: Using the complete class functions encapsulated in step 3, build the Flask server, including loading the RKLLM model based on user-specified parameters at Flask startup, defining the format of user input and the method for parsing the input, differentiating between streaming/non-streaming inference call methods, defining the return format for inference output and the specific implementation of callback functions, and handling RKLLM release.

The specific implementations of the above modules form the main body of the code in rkllm_server/flask_server.py, thereby completing the deployment of the RKLLM-Server-Flask example. Users can modify the initialization definitions for the RKLLM model to implement different custom models. Additionally, users can refer to the RKLLM-Server-Flask deployment example for implementing their own custom server deployment.

3.4.1.3 Client-side: API Access Example

RKLLM-Server-Flask exposes OpenAI API-compatible HTTP endpoints. Client examples are provided for both standard chat completion and Function Calling (tool calling). For more details, please refer to the README.md in the example/rkllm_server_demo directory. Users may also access the service using any OpenAI API-compatible third-party client or application.

3.4.2 Deployment Example of RKLLM-Server-Gradio

Gradio is a simple and easy-to-use Python library used for quickly building interactive interfaces for

machine learning models. In this section, we will specifically introduce how to quickly deploy RKLLM-Server-Gradio on a Linux device using Gradio, and directly access the server within the local network to perform RKLLM model inference. The following Figure shows an example of the web interface after successfully deploying RKLLM-Server-Gradio:




Figure 3-4 External access page for RKLLM-Server deployment example

In the current `rkllm_server_demo` directory, the specific implementation code for the RKLLM-Server-Gradio deployment example shown in Figure 3-3 is included. Users can also directly use the one-click deployment script `build_rkllm_server_gradio.sh` to quickly set up RKLLM-Server-Gradio. After successful deployment, users can choose to access the RKLLM model for inference either through the web interface or via API access.

3.4.2.1 Server-side: Introductions for RKLLM-Server-Gradio Example

The one-click deployment script `build_rkllm_server_gradio.sh` is designed to facilitate the quick setup of RKLLM-Server-Gradio on a Linux development board. Similar to the setup process for RKLLM-Server-Flask, users should pay attention to the following points before using the one-click deployment script:

- 1) Ensure that the development board is connected to the network via an Ethernet cable. Use the `ifconfig` command in the adb shell to query the specific IP address of the development board. RKLLM-Server-Gradio will be set up as a server within the local network using this IP address to accept client

access.

2) Users need to complete the smooth conversion of the RKLLM model and have pushed the RKLLM model to the Linux board before executing the one-click deployment script.

The usage of the one-click deployment script `build_rkllm_server_gradio.sh` is similar to that of `build_rkllm_server_flask.sh`:

```
./build_rkllm_server_gradio.sh
  --workshop [RKLLM-Server Working Path]
  --model_path [Absolute Path of Converted RKLLM Model on Board] --
platform [Target Platform: RK3588/RK3576/RK3562/RV1126B]
  [--lora_model_path [LoRA Model Path]]
  [--prompt_cache_path [Prompt Cache File Path]]
```

Similarly, the `workshop` parameter specifies the working directory for RKLLM-Server-Gradio on the device; the `model_path` parameter indicates the absolute path of the RKLLM model on the device, which was converted using RKLLM-Toolkit, and RKLLM-Server-Gradio will read the model from this path during operation; the `platform` parameter specifies the platform type being used, such as RK3588, RK3576, RK3562 or RV1126B; `lora_model_path` and `prompt_cache_path` are optional parameters, allowing the user to specify the file paths for loading a LoRA model or utilizing the Prompt feature if needed.

Users can directly call the `build_rkllm_server_gradio.sh` script on the PC side (not on the development board) using the following command to quickly deploy the RKLLM-Server-Gradio example:

```
./build_rkllm_server_gradio.sh
  --workshop /user/data
  --model_path /user/data/model.rkllm
  --platform rk3588 --adb_device [ADB Device Serial]
```

After executing the above command, the one-click deployment script will perform the following steps:

- 1) Check the Linux environment on the board.
- 2) Automatically install the Gradio library if not already installed.
- 3) Push the necessary files under `rkllm_server_demo/rkllm_server` to the board.
- 4) Index the RKLLM model in the preset working directory for RKLLM-Server-Gradio.

Once you see the message "RKLLM Model has been initialized successfully!" in the terminal, it

indicates that the RKLLM-Server-Gradio example has been successfully launched.

```
warnings.warn(

=====init....=====
rkllm-runtime version: 1.0.2b9, rknpu driver version: 0.9.7, platform: RK3588
load prompt cache from '/data/cw/prompt_cache.bin'
loaded a prompt cache with prompt size of 27 tokens
RKLLM Model has been initialized successfully!
=====
Running on local URL:  http://0.0.0.0:8080

To create a public link, set `share=True` in `launch()`.
```

Figure 3-5 Successful deployment of RKLLM-Server-Gradio in terminal

By referencing the specific code in `build_rkllm_server_gradio.sh`, users can understand the detailed deployment process of the RKLLM-Server-Gradio example. This can help users deploy custom servers more flexibly. It is important to emphasize that in step 3 of `build_rkllm_server_gradio.sh`, the one-click deployment script automatically synchronizes the current version of RKLLM Runtime to `rkllm_server/lib/librkllmrt.so`. This ensures that `gradio_server.py` indexes `librkllmrt.so` when running, and users need to pay attention to the invocation of `librkllmrt.so` when customizing the server.

3.4.2.2 Server-side: Introductions for RKLLM-Server-Gradio Example

The deployment implementation of RKLLM-Server-Gradio is similar to RKLLM-Server-Flask. It also uses the `ctypes` library to directly call the RKLLM Runtime library to perform RKLLM model inference. The specific deployment implementation process can be referenced in Figure below:




Figure 3-6 Overview of the RKLLM-Server-Flask Deployment Implementation Process

The difference is that RKLLM-Server-Gradio chooses to use the gradio library to implement the server setup and complete communication with the client, providing a simple web-based service. This requires specific handling of the Gradio interface in RKLLM-Server-Gradio as follows:

- 1) The Gradio function library provides complete communication for input and output data, so there is no need to define complex input-output implementations for the server via Flask.
- 2) Gradio is a deployment framework based on different control elements. During usage, it is necessary to call various Gradio components to complete the interface design and specify the trigger conditions, function call logic, and the data flow logic between different components for each element.

Users can refer to the main code in `rkllm_server/gradio_server.py` to understand the specific implementation of RKLLM-Server-Gradio, and by modifying the initialization definitions for the RKLLM model within it, they can implement different custom models. Additionally, users can refer to the deployment example of the RKLLM-Server-Gradio to deploy their own custom server.


3.4.2.3 Client: RKLLM-Server-Gradio Usage Instructions

After successfully deploying the RKLLM-Server-Gradio on a Linux development board, users can access it via two methods: "Interface Access" and "API Access".

1) Interface Access: Upon successfully starting the RKLLM-Server-Gradio with the one-click deployment script, users can directly access the RKLLM model for quick interaction by opening any web browser on a computer within the current local area network and navigating to "board_IP:8080" (e.g., "x.x.x.x:8080" as shown in Figure 3-2). Gradio automatically integrates Markdown, HTML, and other syntaxes, adapting to the format of the RKLLM model's output results, such as code snippets and Markdown text. Additionally, during the setup of RKLLM-Server, an access queue is initiated. When multiple users interact with the RKLLM-Server simultaneously, the inputs are processed and returned in the order they were submitted. It's important to note that when a user's interaction with the RKLLM-Server is in the inference state (i.e., the dialogue box is highlighted), the server will not accept the user's next input until the current inference is completed.

2) API Access: In the rkllm_server_demo directory, chat_api_gradio.py is provided. After installing gradio_client on the PC (using the command: pip install gradio_client), users can interact with the RKLLM-Server solely through the API interface without relying on the graphical interface, as shown in following Figure. Before using chat_api_gradio.py, it's important to modify the IP address in the code to match the current IP address of the development board, as shown in the following code.

```
from gradio_client import Client
client = Client("http://x.x.x.xx:8080/")
```



```
(base) /rkllm/rkllm-runtime/examples/rkllm_server_demo$ python chat_api.py
在终端中输入您的问题，即可与 RKLLM 模型进行对话....
=====
请输入您的问题：请给出福州的介绍
Loaded as API: http://172.16.10.102:8080/ ✓
Q: 请给出福州的介绍
A: 福州，福建省省会，位于中国东南沿海的大陆架上，东临太平洋，西隔台湾海峡与台湾岛相望；北濒东海，南界巴士海峡与菲律宾群岛相对。全市总面积13247平方千米，其中市区面积6580平方千米。福州市区气候温暖湿润，四季分明，是著名的“鱼米之乡”。福州是中国东南沿海重要的港口城市和风景名胜区，拥有众多的历史文化遗产和自然景观。
请输入您的问题： █
```

Figure 3-7 Access the RKLLM-Server-Gradio via API calls in terminal

Users can choose between the two client invocation methods based on their specific needs. For instance, when providing interactive services within a local area network, it's recommended to use the interface access method. On the other hand, if customizing access behaviors to RKLLM-Server-Gradio is required, it's advisable to use API Access for further development.

Lastly, it's important to note that in the implementation of RKLLM-Server-Gradio, there isn't a definition of data structures for sending and receiving data similar to OpenAI-API. Therefore, this deployment implementation is not compatible with the OpenAI-API interface. When conducting further development, users should refer to the specific function implementation in chat_api_gradio.py. If compatibility with the OpenAI-API interface is required, please refer to the implementation of RKLLM-Server-Flask.

4 Reference

RKLLM: <https://github.com/airockchip/rknn-llm>

RKNN: <https://github.com/airockchip/rknn-toolkit2>

Rockchip